Homework 2: MapReduce Concepts & Spark Fundamentals

Points: 20 | **Due:** Sunday, February 8, 2026 @ 11pm Pacific

Author: Richard Young, Ph.D. | UNLV Lee Business School

Compute: CPU (free tier is fine)

Learning Objectives

1. **Set up** Apache Spark on Google Colab
2. **Understand** how Spark partitions data for parallel processing
3. **Measure and compare** processing performance with different configurations
4. **Apply** K-Means clustering and interpret business results
5. **Explain** distributed computing through your own diagram

Grading

Component	Points	What We're Looking For
Spark Setup	3	Create a Spark session and explain configuration
Data Partitioning	5	Demonstrate understanding of data distribution
Performance Experiment	5	Analyze why speedup isn't linear
K-Means Clustering	5	Apply clustering and interpret results
Diagram	2	Clear diagram explaining distributed processing
Total	**20**	

Instructions

1. Open `MIS769_HW2_MapReduce_Spark.ipynb` in Google Colab
2. Complete Part 1: Install Spark and create a session
3. Complete Part 2: Explore data partitioning
4. Complete Part 3: Run performance experiment with 1, 2, and 4 cores
5. Complete Part 4: Apply K-Means to Netflix data and name your clusters
6. Complete Part 5: Draw your own diagram of how Spark processes data

Questions to Answer

- **Q1:** What does `local[2]` mean? What would `local[4]` do differently?
- **Q2:** Why doesn't 4 cores give exactly 4x speedup?
- **Q3:** What characterizes each cluster? Give them descriptive names.

Submission

Upload to Canvas: - Your completed `.ipynb` notebook with all cells executed

---

— Richard Young, Ph.D.