

Homework 3: Text Processing Fundamentals

Points: 20 | **Due:** Sunday, February 15, 2026 @ 11pm Pacific

Author: Richard Young, Ph.D. | UNLV Lee Business School

Compute: CPU (free tier)

Learning Objectives

1. **Install and use** NLP libraries (spaCy, NLTK)
 2. **Understand WHY** we preprocess text (not just how)
 3. **Create** domain-specific stopwords for your data
 4. **Identify** cases where cleaning hurts your analysis
-

Why This Matters for Business

Market Research: Before analyzing thousands of customer reviews, companies like Procter & Gamble preprocess text to focus on meaningful words. Poor preprocessing = misleading insights = bad product decisions.

Sentiment Analysis: When United Airlines analyzes social media mentions, removing “not” from “not satisfied” completely inverts the meaning. Smart text processing preserves business-critical signals.

Search & Discovery: E-commerce platforms like Shopify tune their stopword lists per industry. A wine retailer keeps “dry” (meaningful); a weather app removes it (noise).

Grading

Component	Points	Effort	What We're Looking For
Environment Setup	3	*	NLP libraries installed and working
Standard Stopwords	4	*	Applied and analyzed impact
Domain Stopwords	5	**	Created 10+ with justification
Negation Analysis	5	**	Smart stopwords preserving meaning
Visualization	3	*	Word cloud comparison
Total	20		

Effort Key: * Straightforward | ** Requires thinking | *** Challenge

The Big Picture

Text preprocessing is like preparing ingredients before cooking. But just like cooking, **the same preparation isn't right for every dish.**

Standard stopword removal can destroy important meaning (like negations in sentiment analysis).




Figure 1: Text Tokenization Pipeline

Instructions

1. Open MIS769_HW3_Text_Processing.ipynb in Google Colab
2. Load your dataset (HuggingFace, Kaggle, or your own)
3. Apply standard stopword removal and analyze what was removed
4. Create your own domain-specific stopwords with justification
5. Analyze the negation problem and create a “smart” stopword list
6. Generate word cloud visualizations comparing approaches

What Your Output Should Look Like

Stopword Removal Impact:

STOPWORD REMOVAL IMPACT

=====

Average reduction: 45.2%

Median reduction: 44.8%

Top Words After Cleaning:

TOP 30 MOST COMMON WORDS (after standard cleaning)

1. movie	12,847
2. film	9,234
3. good	7,892
4. like	6,543
...	

Negation Problem Example:

Original: This product is not good at all
Cleaned: product good ← MEANING REVERSED!

Common Mistakes (and How to Avoid Them)

Mistake	Symptom	Fix
Forgetting to lowercase	“Movie” and “movie” counted separately	Add <code>.lower()</code> before tokenizing
Removing all punctuation first	Can’t detect sentence boundaries	Keep punctuation until after sentence splitting
Using wrong stopword list	British vs American English issues	Check: ‘colour’ in stopwords
Not restarting runtime	<code>ModuleNotFoundError</code> after pip install	Runtime □ Restart runtime
Treating numbers as stopwords	Lose important data like prices, ratings	Keep numbers if relevant to your domain

If you see this error:

`LookupError`: Resource stopwords `not` found.

Run: `nltk.download('stopwords')`

Questions to Answer

- **Q1:** Which removed stopwords might carry meaning in your domain?
 - **Q2:** Why did you choose each domain-specific stopword?
 - **Q3:** When should you preserve vs. remove negations?
 - **Q4:** Give an example where aggressive text cleaning would HURT your analysis results.
-

Going Deeper (Optional Challenges)

Challenge A: Lemmatization Comparison

Compare results using stemming (Porter) vs. lemmatization (spaCy). Which produces more interpretable results for business reporting?

Challenge B: Multi-language Stopwords

If your dataset contains non-English text, implement stopword removal for a second language. Compare the challenges.

Challenge C: Context-Aware Stopwords

Build a function that removes stopwords only when they appear in certain contexts (e.g., keep “not” before adjectives, remove it elsewhere).

Quick Reference

```
# Load libraries
import nltk
from nltk.corpus import stopwords
import spacy
nlp = spacy.load("en_core_web_sm")

# Get standard stopwords
STOPWORDS = set(stopwords.words('english'))

# Remove stopwords
def clean(text):
    words = text.lower().split()
    return ' '.join([w for w in words if w not in STOPWORDS])

# Preserve negations
NEGATIONS = {'not', 'no', 'never', 'neither', "n't", 'nor'}
SMART_STOPWORDS = STOPWORDS - NEGATIONS

# spaCy tokenization (better than split)
doc = nlp(text)
tokens = [token.text for token in doc]

# Word frequency
from collections import Counter
freq = Counter(all_words).most_common(20)
```

Useful spaCy Token Attributes: | Attribute | Description | Example | | --- | --- | --- |
| token.text | Original text | “running” | | token.lemma_ | Base form | “run” | | token.pos_ | Part of speech | “VERB” | | token.is_stop | Is stopword? | True/False |

Submission

Upload to Canvas: - Your completed .ipynb notebook with all cells executed

— *Richard Young, Ph.D.*