

Homework 3: Text Processing Fundamentals

Points: 20 | Due: See WebCampus for deadline

Author: Richard Young, Ph.D. | UNLV Lee Business School

Compute: CPU (free tier)

Learning Objectives

1. **Install and use** NLP libraries (spaCy, NLTK) in Google Colab
 2. **Understand** WHY we preprocess text (not just how)
 3. **Create** domain-specific stopwords for your data
 4. **Measure** the impact of text cleaning on analysis
 5. **Identify** cases where cleaning hurts your analysis
-

Why This Matters for Business

Market Research: Before analyzing thousands of customer reviews, companies like Procter & Gamble preprocess text to focus on meaningful words. Poor preprocessing = misleading insights = bad product decisions.

Sentiment Analysis: When United Airlines analyzes social media mentions, removing “not” from “not satisfied” completely inverts the meaning. Smart text processing preserves business-critical signals.

Search & Discovery: E-commerce platforms like Shopify tune their stopword lists per industry. A wine retailer keeps “dry” (meaningful); a weather app removes it (noise).

Grading

| Component | Points | Effort | What We're Looking For |
|--------------------|-----------|--------|-------------------------------------|
| Environment Setup | 3 | * | NLP libraries installed and working |
| Standard Stopwords | 4 | * | Applied and analyzed impact |
| Domain Stopwords | 5 | ** | Created 10+ with justification |
| Negation Analysis | 5 | ** | Smart stopwords preserving meaning |
| Visualization | 3 | * | Word cloud comparison |
| Total | 20 | | |

Effort Key: * Straightforward | ** Requires thinking | *** Challenge

The Big Picture

Text preprocessing is like preparing ingredients before cooking. But just like cooking, **the same preparation isn't right for every dish.**

Standard stopword removal can destroy important meaning (like negations in sentiment analysis).

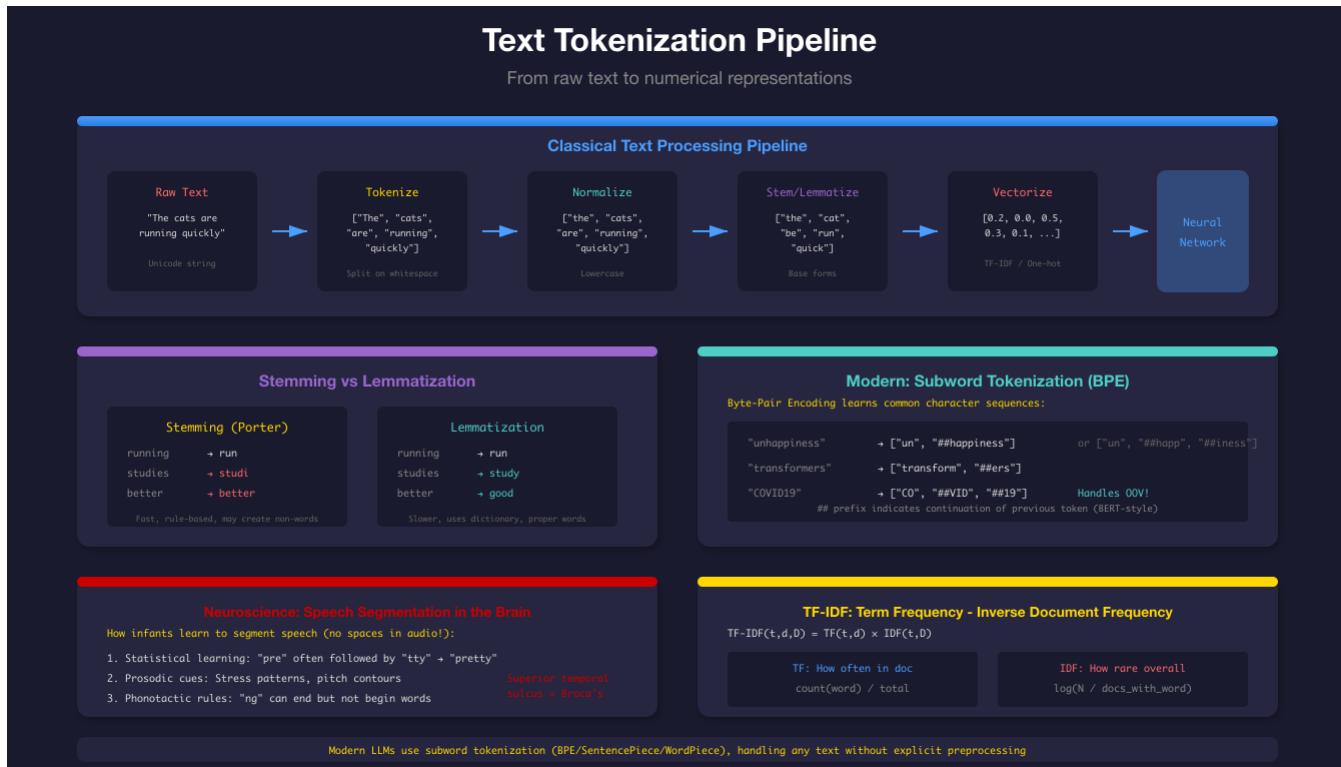


Figure 1: Text Tokenization Pipeline

Instructions

1. Open MIS769_HW3_Text_Processing.ipynb in Google Colab
2. Load your dataset (HuggingFace, Kaggle, or your own)
3. Apply standard stopword removal and analyze what was removed
4. Create your own domain-specific stopwords with justification
5. Analyze the negation problem and create a “smart” stopword list
6. Generate word cloud visualizations comparing approaches

What Your Output Should Look Like

Stopword Removal Impact:

□ STOPWORD REMOVAL IMPACT

Average reduction: 45.2%

Median reduction: 44.8%

Top Words After Cleaning:

- TOP 30 MOST COMMON WORDS (after standard cleaning)

| | |
|----------|--------|
| 1. movie | 12,847 |
| 2. film | 9,234 |
| 3. good | 7,892 |
| 4. like | 6,543 |
| ... | |

Negation Problem Example:

Original: This product is not good at all
 Cleaned: product good ← MEANING REVERSED!

Common Mistakes (and How to Avoid Them)

| Mistake | Symptom | Fix |
|--------------------------------|--|---|
| Forgetting to lowercase | “Movie” and “movie” counted separately | Add .lower() before tokenizing |
| Removing all punctuation first | Can’t detect sentence boundaries | Keep punctuation until after sentence splitting |
| Using wrong stopword list | British vs American English issues | Check: 'colour' in stopwords |
| Not restarting runtime | ModuleNotFoundError after pip install | Runtime □ Restart runtime |
| Treating numbers as stopwords | Lose important data like prices, ratings | Keep numbers if relevant to your domain |

If you see this error:

LookupError: Resource stopwords **not** found.

Run: `nltk.download('stopwords')`

Questions to Answer

- **Q1:** Which removed stopwords might carry meaning in your domain?
 - **Q2:** Why did you choose each domain-specific stopword?
 - **Q3:** When should you preserve vs. remove negations?
-

Submission



Richard Young, Ph.D.

Going Deeper (Optional Challenges)

Challenge A: Lemmatization Comparison [10%]

Compare results using stemming (Porter) vs. lemmatization (spaCy). Which produces more interpretable results for business reporting?

Challenge B: Multi-language Stopwords

If your dataset contains non-English text, implement stopword removal for a second language. Compare the challenges.

Challenge C: Context-Aware Stopwords

Build a function that removes stopwords only when they appear in certain contexts (e.g., keep “not” before adjectives, remove it elsewhere).

Quick Reference

```
# Load libraries
import nltk
from nltk.corpus import stopwords
import spacy
nlp = spacy.load("en_core_web_sm")

# Get standard stopwords
STOPWORDS = set(stopwords.words('english'))

# Remove stopwords
def clean(text):
    words = text.lower().split()
    return ' '.join([w for w in words if w not in STOPWORDS])

# Preserve negations
NEGATIONS = {'not', 'no', 'never', 'neither', "n't", 'nor'}
SMART_STOPWORDS = STOPWORDS - NEGATIONS

# spaCy tokenization (better than split)
doc = nlp(text)
tokens = [token.text for token in doc]

# Word frequency
from collections import Counter
freq = Counter(all_words).most_common(20)
```

Useful spaCy Token Attributes: | Attribute | Description | Example | |-----|-----|-----|
 | token.text | Original text | “running” | | token.lemma_ | Base form | “run” | | token.pos_ | Part of speech
 | “VERB” | | token.is_stop? | Is stopword? | True/False |
