

Homework 5: Named Entity Recognition (NER)

Points: 20 | **Due:** Sunday, March 1, 2026 @ 11pm Pacific

Author: Richard Young, Ph.D. | UNLV Lee Business School

Compute: CPU (free tier) — GPU recommended for bonus

Learning Objectives

1. **Understand** Named Entity Recognition concepts
 2. **Use** spaCy for entity extraction
 3. **Extract** business-relevant entities from text data
 4. **Analyze** entity patterns in reviews
 5. **Visualize** entity distributions
-

Why This Matters for Business

Brand Monitoring: Companies like Coca-Cola use NER to automatically track mentions of their brands, competitors, and products across millions of social media posts daily. Manual tracking at this scale is impossible.

Financial Analysis: Bloomberg's terminal uses NER to extract company names, executives, and monetary values from news articles—feeding trading algorithms and analyst dashboards in real-time.

Legal Discovery: Law firms use NER to scan thousands of documents for people, organizations, and dates relevant to a case. What took paralegals weeks now takes hours.

Supply Chain Intelligence: Walmart monitors news for location entities (GPE) to predict disruptions—a hurricane heading toward a supplier's city triggers automatic inventory adjustments.

Grading

Component	Points	Effort	What We're Looking For
Setup & Data	3	*	spaCy loaded, data ready
NER Demo	4	*	Understand entity types
Entity Extraction	6	**	Extract ORG, PRODUCT, GPE from reviews
Analysis	4	**	Identify patterns in entity mentions
Visualization	3	*	Clear bar charts of top entities
Total	20		

Effort Key: * Straightforward | ** Requires thinking | *** Challenge

Entity Types

spaCy recognizes these business-relevant entities:

Entity	Description	Example
ORG	Organizations	“Apple”, “Amazon”
PRODUCT	Products	“iPhone”, “Kindle”
GPE	Countries, cities	“New York”, “USA”
PERSON	People names	“Tim Cook”
MONEY	Monetary values	“\$50”
DATE	Dates and periods	“September 2024”
EVENT	Named events	“Olympics”, “Black Friday”

Instructions

1. Open MIS769_HW5_NER.ipynb in Google Colab
2. Load spaCy and understand entity types
3. Extract entities from your review dataset
4. Analyze: Which organizations/products are mentioned most?
5. Visualize the top entities by category

What Your Output Should Look Like

NER Demo Output:

NAMED ENTITIES FOUND

Apple		ORG		Companies, agencies, institutions
Tim Cook		PERSON		People, including fictional
iPhone 15		PRODUCT		Objects, vehicles, foods, etc.
San Francisco		GPE		Countries, cities, states
\$999		MONEY		Monetary values
September 22, 2024		DATE		Absolute or relative dates

Entity Counts:

- Extraction complete!
- ORG: 8,234
- PRODUCT: 1,892
- GPE: 4,567
- PERSON: 12,453

Top Organizations:

TOP ORGANIZATIONS

Hollywood		1,247
BBC		523
HBO		412

Disney		389
Netflix		298

Common Mistakes (and How to Avoid Them)

Mistake	Symptom	Fix
Not downloading spaCy model	OSError: Can't find model	Run !python -m spacy download en_core_web_sm
Processing very long texts	Slow/crashes	Truncate: text[:5000]
Expecting perfect accuracy	Frustration	NER is ~85% accurate; document errors
Forgetting to cast to string Case sensitivity confusion	TypeError on null values “apple” (fruit) vs “Apple” (company)	Use str(text) NER uses context, not just capitalization

If you see this error:

OSError: [E050] Can't find model 'en_core_web_sm'

Run:

!python -m spacy download en_core_web_sm

Then restart runtime (Runtime □ Restart runtime)

If extraction is too slow: - Process a sample: df.sample(1000) - Truncate long texts: text[:5000]
- Use nlp.pipe() for batch processing (faster)

Questions to Answer

- **Q1:** Were the extracted entities accurate? What errors did you observe?
 - **Q2:** What business insights can you derive from entity mentions?
 - **Q3:** How could you improve NER for your specific domain?
-

Going Deeper (Optional Challenges)

Challenge A: Entity Co-occurrence Network

Build a network graph showing which entities appear together. Do certain companies get mentioned alongside certain people? Use networkx to visualize.

Challenge B: Sentiment by Entity

For each organization mentioned, calculate the average sentiment of reviews that mention it. Which companies have positive vs. negative associations?

Challenge C: Custom Entity Training

Train a custom NER model to recognize domain-specific entities not in spaCy's default model (e.g., FEATURE, COMPETITOR, PRODUCT_LINE). Requires the en_core_web_trf model and GPU.

Quick Reference

```
# Load spaCy
import spacy
nlp = spacy.load("en_core_web_sm")

# Process text
doc = nlp("Apple CEO Tim Cook announced the new iPhone in San Francisco.")

# Extract entities
for ent in doc.ents:
    print(f"{ent.text[:20]} | {ent.label_[:10]} | {spacy.explain(ent.label_)}")

# Get specific entity types
orgs = [ent.text for ent in doc.ents if ent.label_ == "ORG"]
people = [ent.text for ent in doc.ents if ent.label_ == "PERSON"]

# Batch processing (faster for large datasets)
texts = ["Text 1", "Text 2", "Text 3"]
for doc in nlp.pipe(texts, batch_size=50):
    entities = [(ent.text, ent.label_) for ent in doc.ents]

# Count entities
from collections import Counter
all_orgs = []
for doc in nlp.pipe(texts):
    all_orgs.extend([ent.text for ent in doc.ents if ent.label_ == "ORG"])
org_counts = Counter(all_orgs).most_common(10)

# Visualize entities in notebook
from spacy import displacy
displacy.render(doc, style="ent", jupyter=True)
```

Entity Labels Cheat Sheet: | Label | Meaning | Business Use | |---|---|---|---| | ORG | Organization | Competitor tracking | | PERSON | Person name | Executive mentions | | GPE | Geo-political entity | Market analysis | | PRODUCT | Product name | Product mentions | | MONEY | Currency amounts | Pricing intelligence | | DATE | Date/time | Timeline analysis | | EVENT | Named event | Event impact | | NORP | Nationality/group | Demographic insights |

Submission

Upload to Canvas: - Your completed .ipynb notebook with all cells executed

— *Richard Young, Ph.D.*